

FPGA implementations of Histograms of Oriented Gradients in FPGA

C. Bourrasset¹, L. Maggiani^{2,3}, C. Salvadori^{2,3}, J. Sérot¹, P. Pagano^{2,3}
and F. Berry¹



Scuola Superiore
Sant'Anna

¹ Institut Pascal- D.R.E.A.M - Aubière, France



consorzio nazionale
interuniversitario
per le telecomunicazioni

² TeCIP Institute - Scuola Superiore Sant'Anna - Pisa

³ CNIT – National Laboratory of Photonic Networks - Pisa



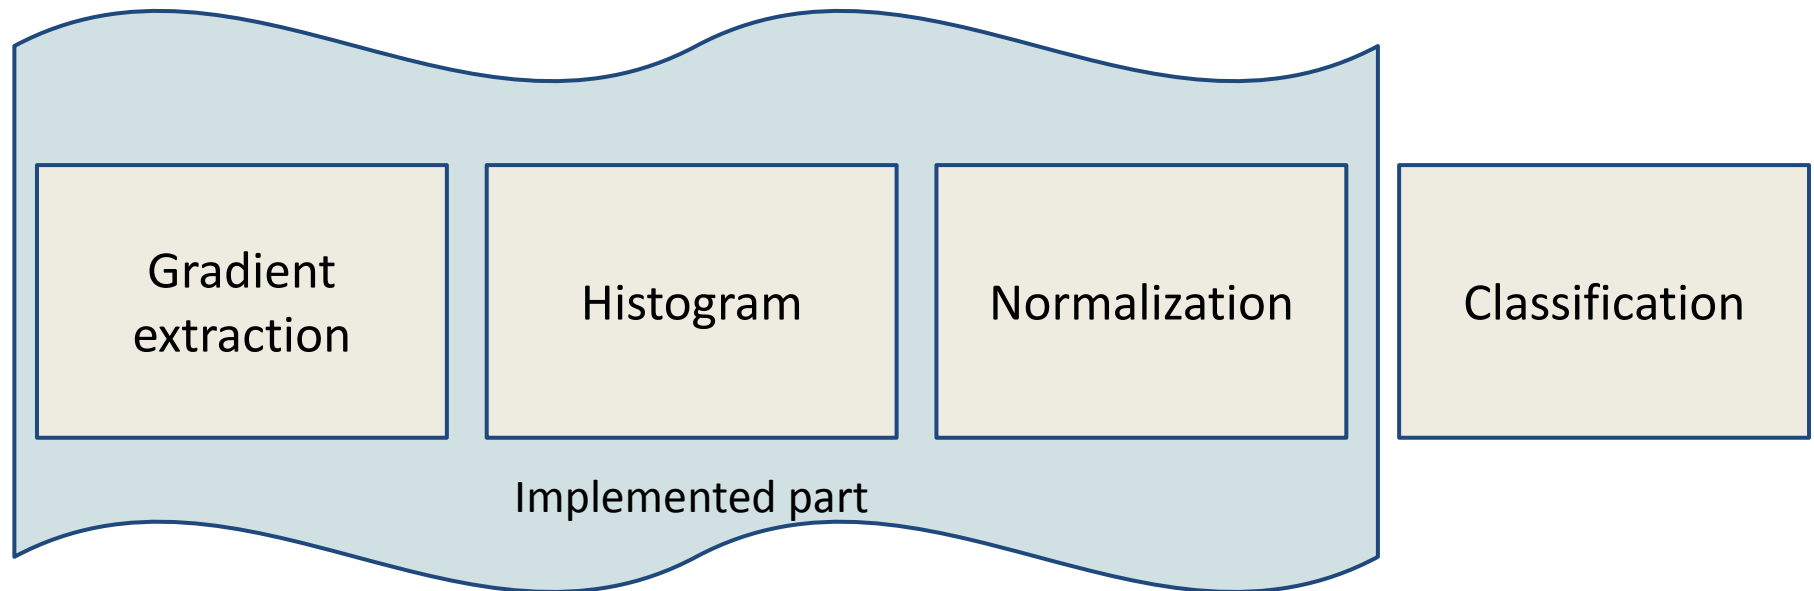
Outline

- **Introduction**
 - Histogram of Oriented Gradients (HOG) pipeline
 - Hardware vs. Software implementation
- **Hardware-Software Codesign approach**
 - Gradient extraction
 - Histogram generation
 - Normalization
 - Implementation results
- **Domain Specific Language approach**
 - CAPH language
 - Implementation results

Introduction

HOG pipeline¹

In 2005 Dalal presents the HOG pipeline



¹ *Histograms of Oriented Gradients for Human Detection*, Dalal and Triggs, INRIA, 2005

Gradient Extraction

$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$$

$$\frac{\partial I}{\partial x} \equiv I(x+1, y) - I(x-1, y)$$

$$\frac{\partial I}{\partial y} \equiv I(x, y+1) - I(x, y-1)$$

Gradient $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}$

■ $f_x =$

-1	0	+1
----	---	----

■ $f_y =$

-1
0
+1





Histogram and normalisation

- The matrix of gradient is divided in cells
- Histogram of gradient orientations for each cell...
 - ... weighted by the gradient magnitude
- The adjacent cells are grouped in blocks
- The cell histograms inside are normalised...
 - .. in order to equalise the luminance among close cells


cell

block



Support vector machine (SVM)

- SVM is [supervised learning](#) binary classifier
- Given a set of training examples, an SVM training algorithm builds a model that assigns new examples into one category or the other
- New examples are then mapped into that same space and predicted to belong to a category based on which side of the gap they fall on²



² Wikipedia



FPGA implementations of Histograms of Oriented Gradients for vehicle detection

The Hardware-Software codesign approach

Workshop on Architecture of Smart Camera, Sevilla, June 2013

Luca Maggiani^(*), Claudio Salvadori^(*), Paolo Pagano^(**)

^(*) Tecip Institute - Scuola Superiore Sant'Anna - Pisa


^(**) CNIT – National Laboratory of Photonic Networks - Pisa

Why we need an optimized solution?

Code oriented approach

NIOSII CPU, 6-stage pipeline, 50MHz

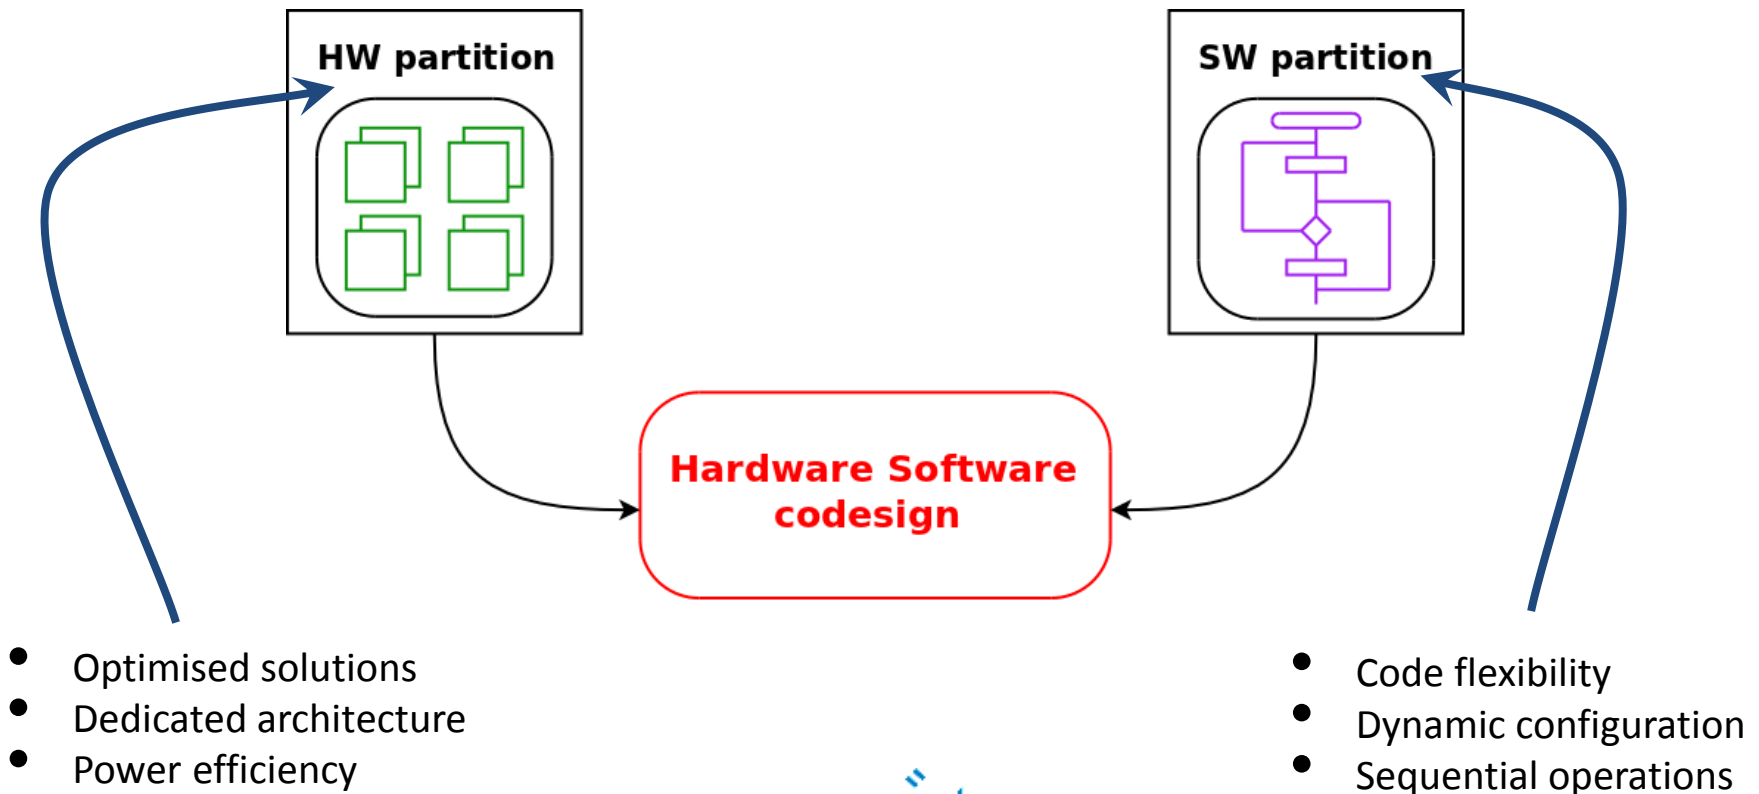
320x240, YUV422, CMOS Camera OV7670 @15fps



Platform	Function	nclk * 10 ⁶	T (ms)	fps
SW elaboration NiosII/f	Camera + LCD	3.965	79.3	12.6 over 15
SW elaboration NiosII/f	Edge detector	37.700	754	1.32 over 15
SW elaboration NiosII/e	Edge detector	77.900	1550	0.64 over 15
SW elaboration NiosII/f	Backgrd Sub.	13.951	279	3.58 over 15

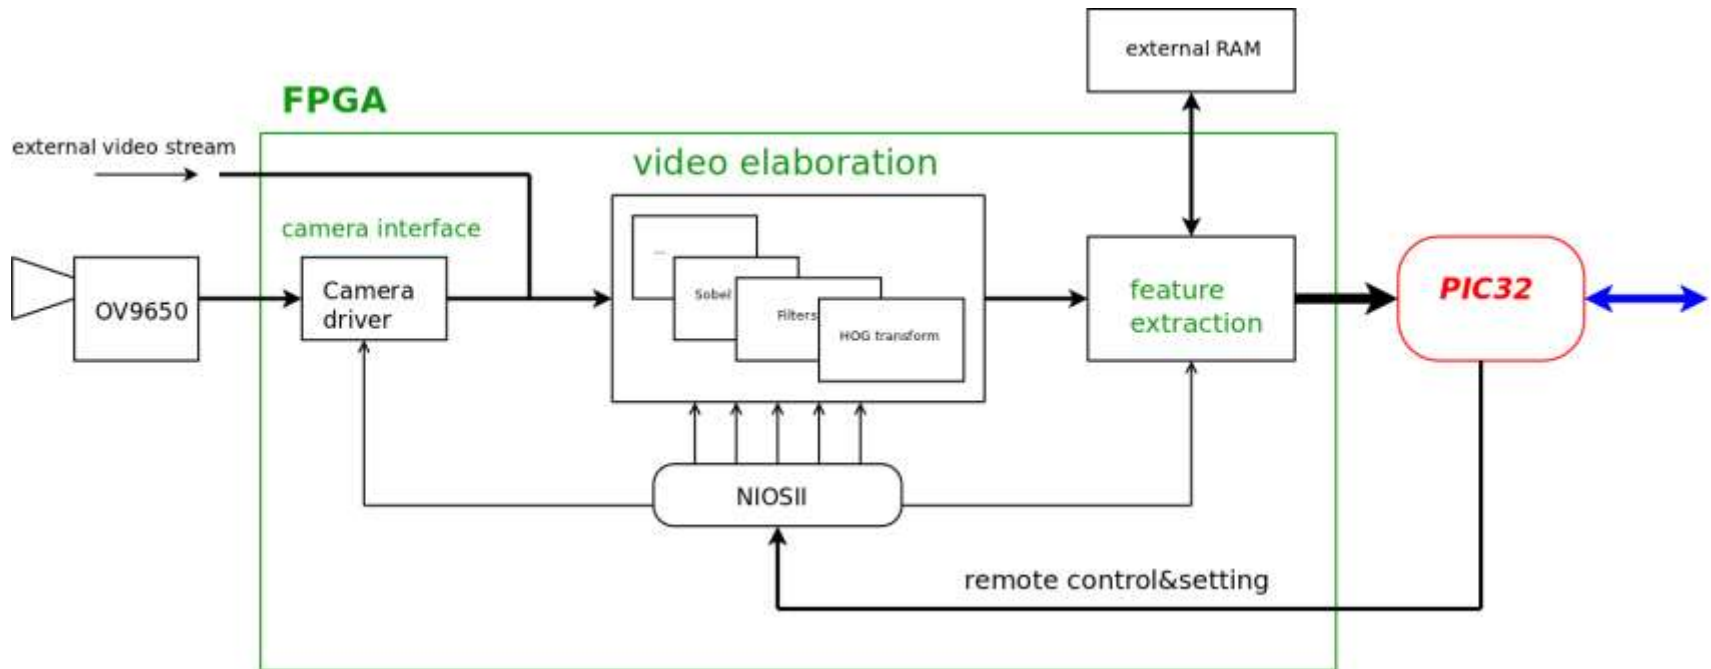
Hardware software codesign

Hardware-Software codesign: Hardware and Software joint development technique to exploit both the HDL optimisation and the code flexibility



HOG pipeline: FPGA implementation

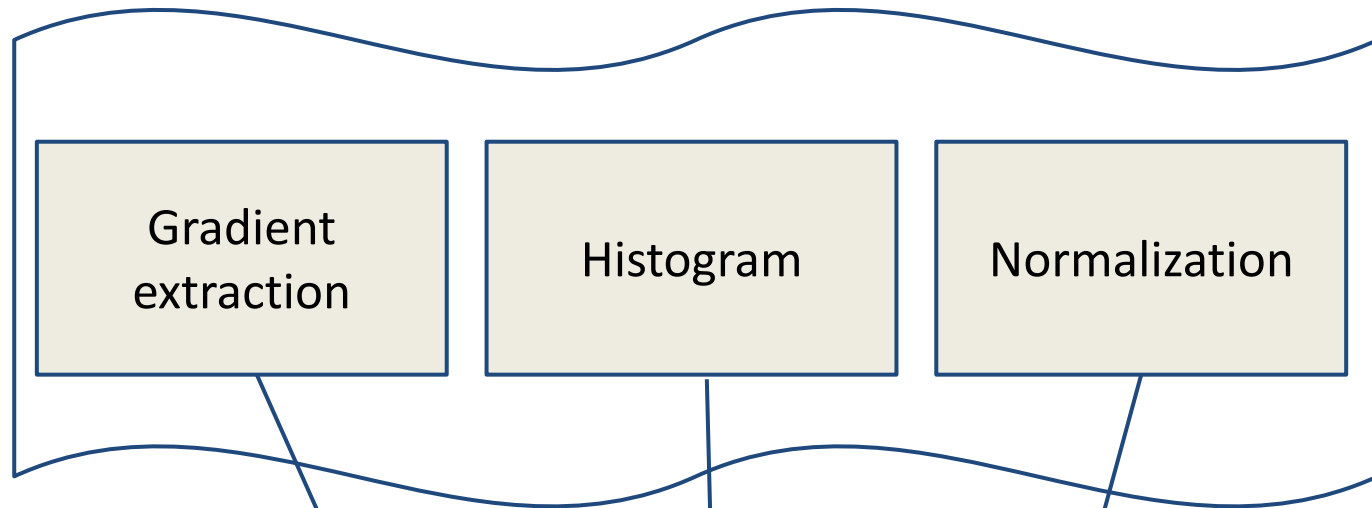
HOG algorithm implementation



Each hardware block is implemented using the **streaming paradigm**:


- data are processed when appear in input

FPGA implementation of the HOG pipeline



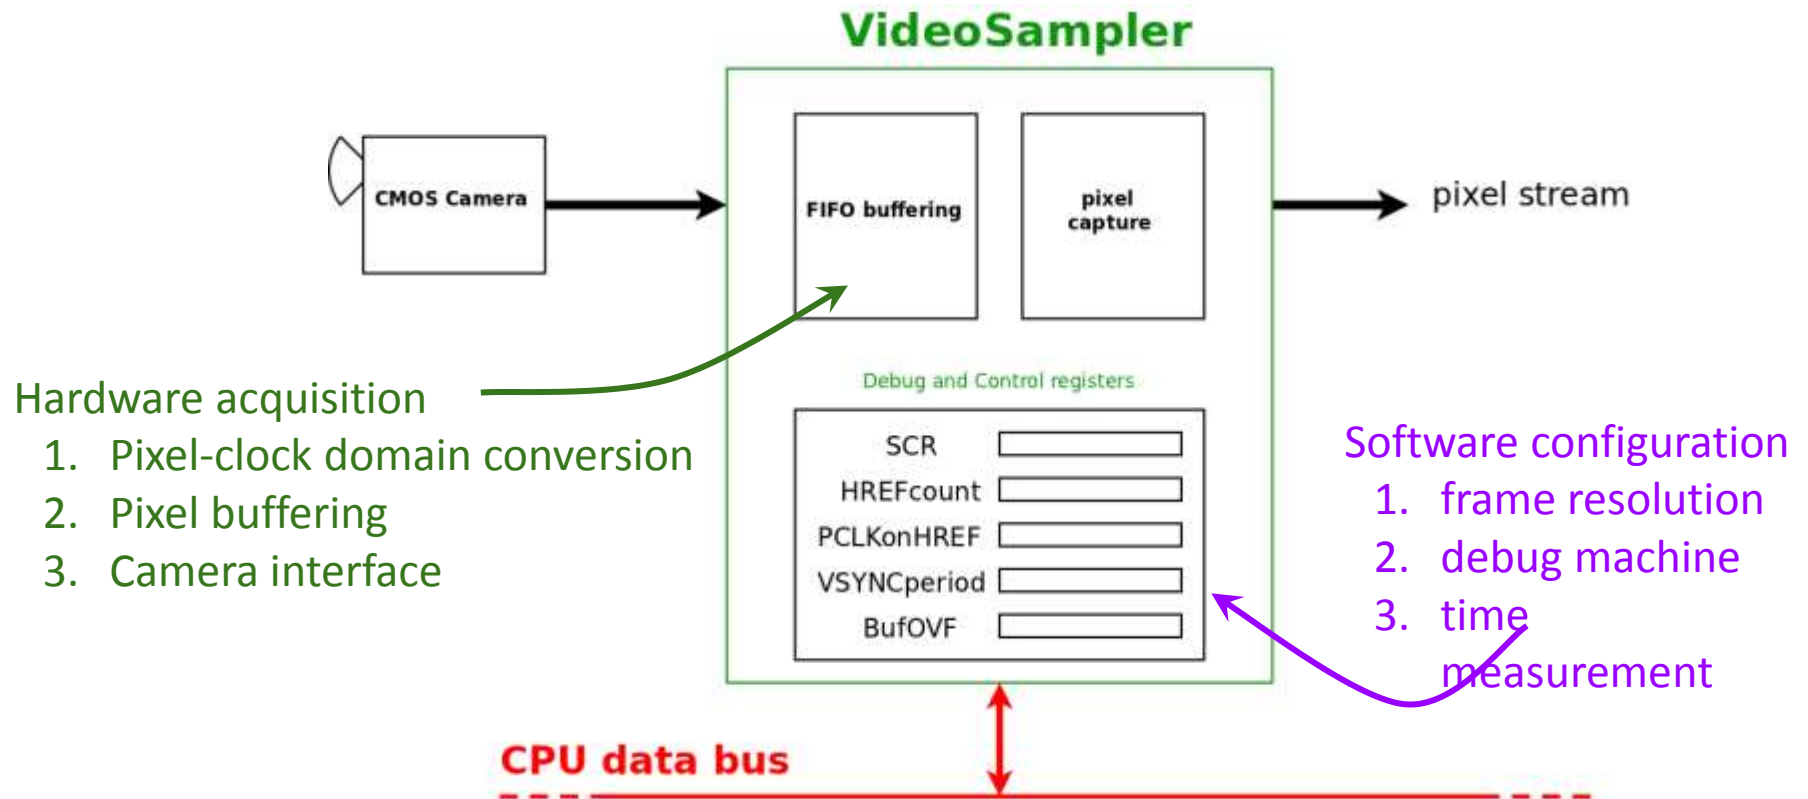
FPGA implementation

FPGA

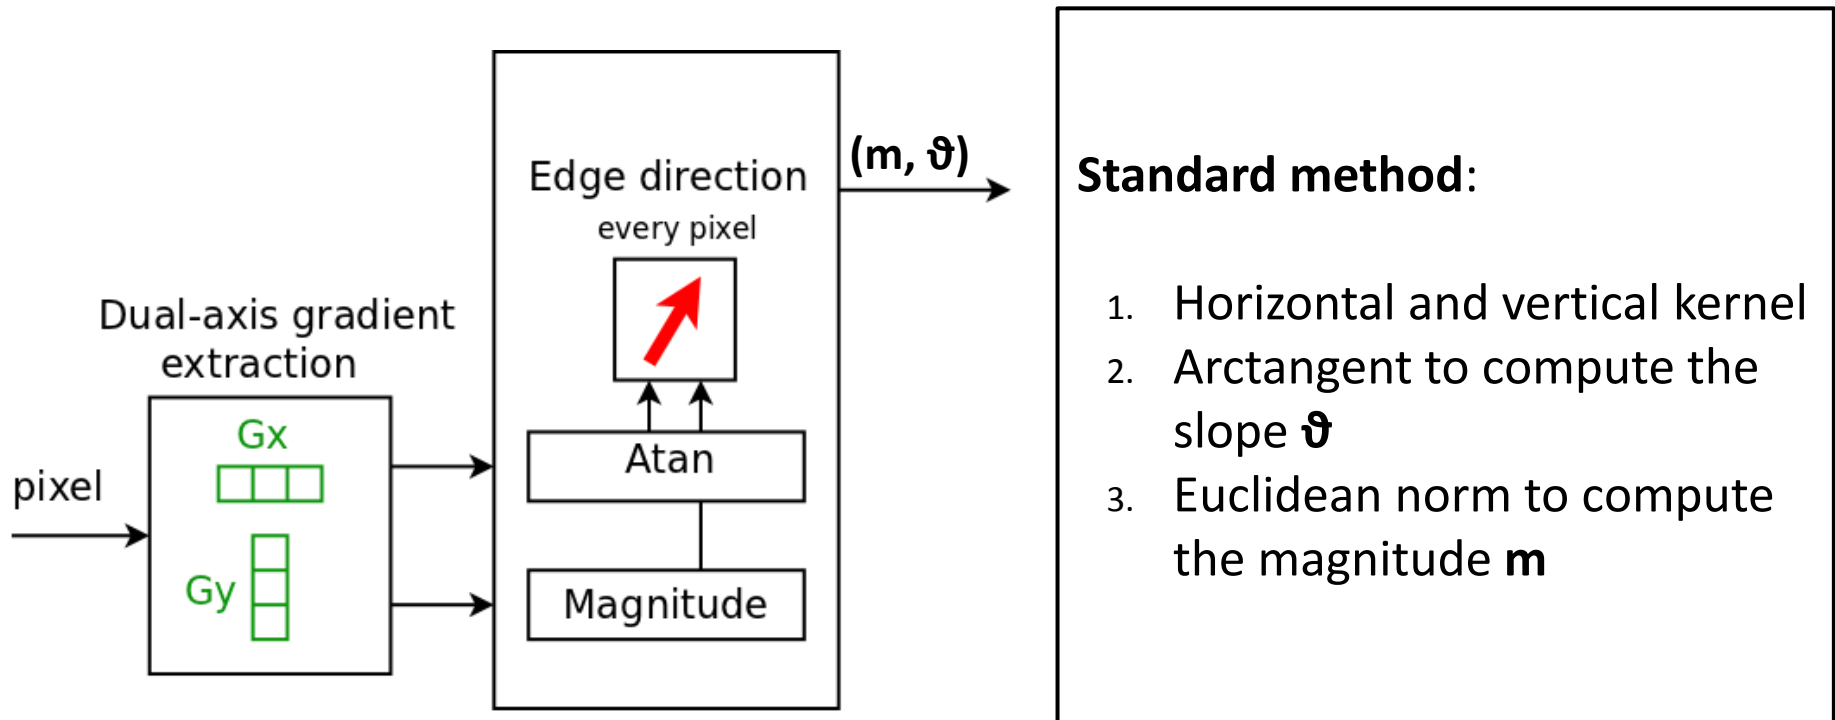


Video Sampler

Performs the video acquisition from a CMOS Camera



Gradient extraction



To implement on HW both "atan" and "sqrt", the usage of a Look-Up Table (LUT) is needed:

- massive memory access (the processing latency is increased)

Hardware gradient extraction (1/3)

Mathematical point of view:

1. The corner domain is sampled in N bins called ϑ_k
 - a. The interval $[0, \pi]$ is considered, as in the original HOG algorithm
2. The idea is to approximate the complete set of versors using a subset of N elements:



$$\hat{I} = \{\hat{i} = \hat{x}\cos\theta_k + \hat{y}\sin\theta_k | k = 0 \dots N-1\}$$

1. ρ and ϑ are computed as:

$$r = \arg\{MAX\{\bar{v} \cdot \hat{i}_k | \hat{i}_k \in \hat{I}\}\}$$

$$\theta \simeq \theta_r$$

$$\rho \simeq \bar{v} \cdot \hat{i}_r$$



Hardware gradient extraction (2/3)

Practically:

1. starting from the **spatial gradient definition**...
1. It is possible to compute its component along the previously introduced versor \hat{i}_k
1. Then, expliciting the spatial gradient and carrying out the **dot-product**, we obtain the matrix coefficients (**edge detector kernel**)

$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$$

$$\frac{\partial I}{\partial \hat{i}_k} = \nabla I \cdot \hat{i}_k$$

$$\frac{\partial I}{\partial x} \equiv I(x+1, y) - I(x-1, y)$$

$$\frac{\partial I}{\partial y} \equiv I(x, y+1) - I(x, y-1)$$

$$\frac{\partial I}{\partial \hat{i}_k} = \cos\theta \left[I(x+1, y) - I(x-1, y) \right] + \sin\theta \left[I(x, y+1) - I(x, y-1) \right]$$


Hardware gradient extraction (3/3)

Edge detector kernel matrix:

every \hat{i}_k vector generates a 3x3 matrix, which is used as a **kernels** above the image

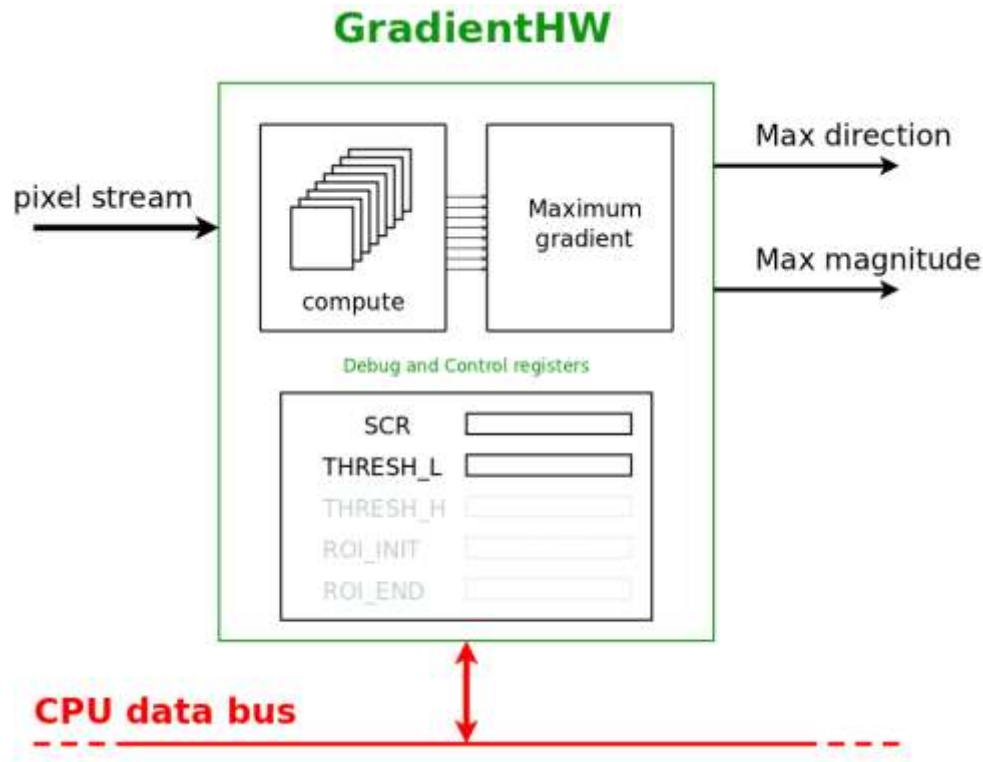
0	$\sin\vartheta_k$	0
$-\cos\vartheta_k$	0	$\cos\vartheta_k$
0	$-\sin\vartheta_k$	0

Exploiting the **hardware parallelism**, we can process a generic N kernels and obtain a Gradient with a resolution $2\pi/N$



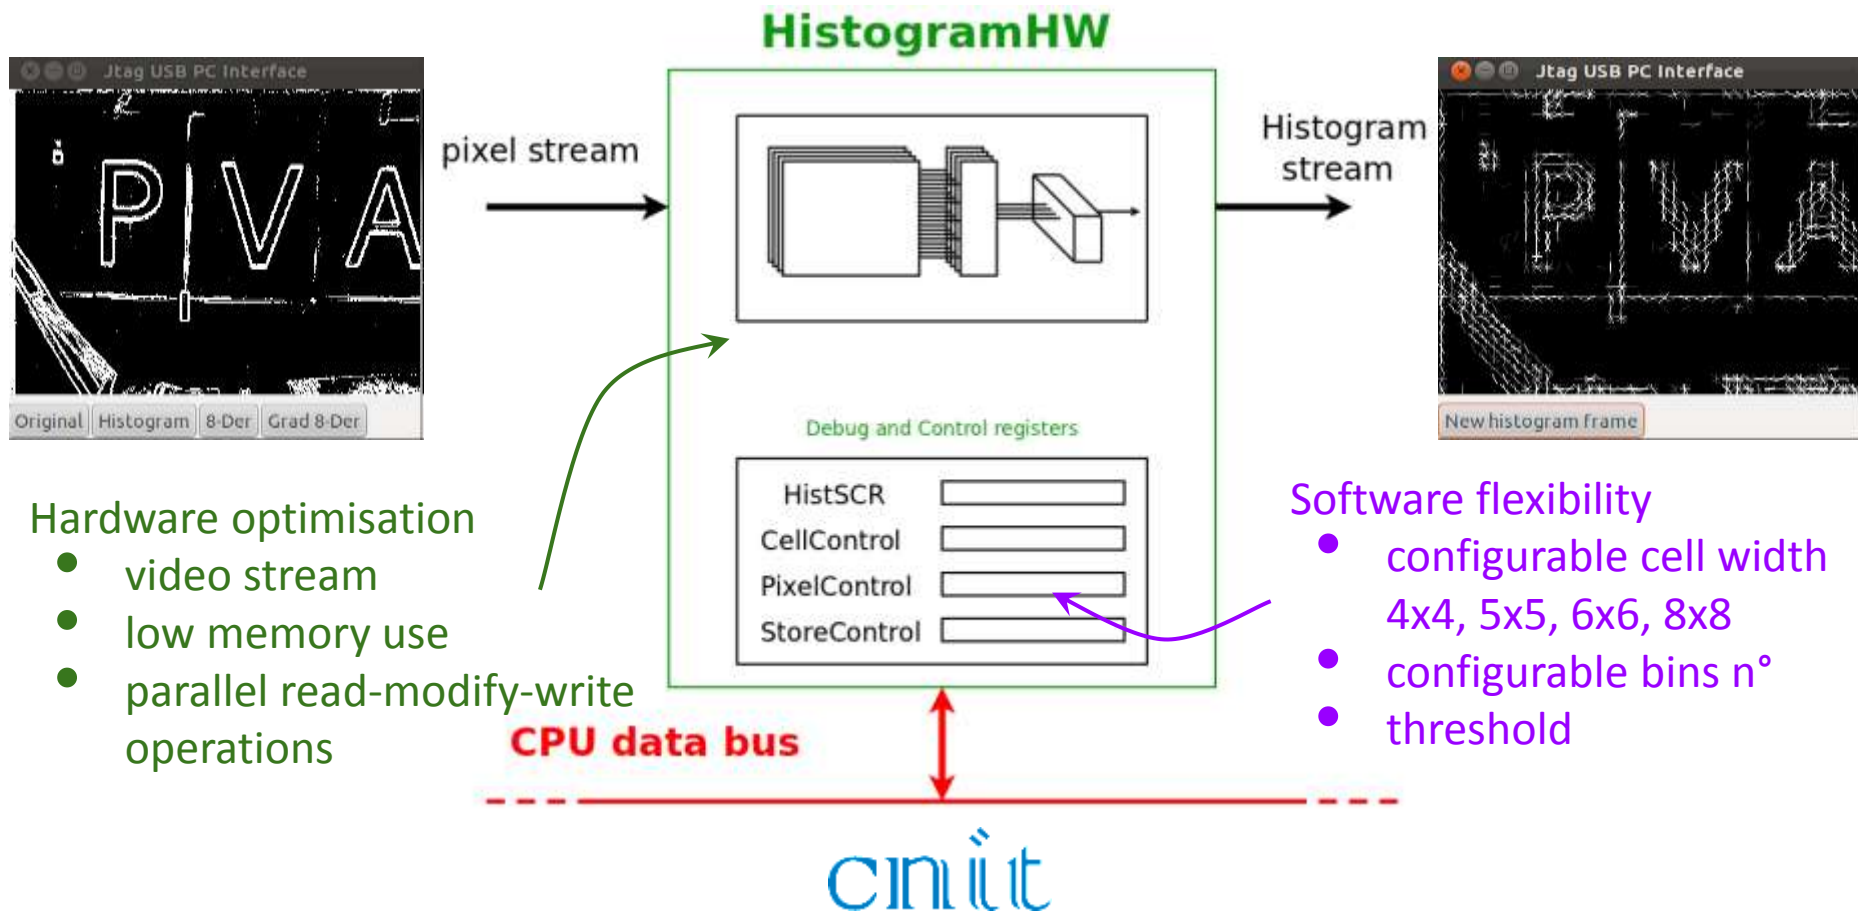
GradientHW

Performs a spatial gradient extraction,
with a **fixed result latency** (2 clock cycle)



HistogramHW (1/3)

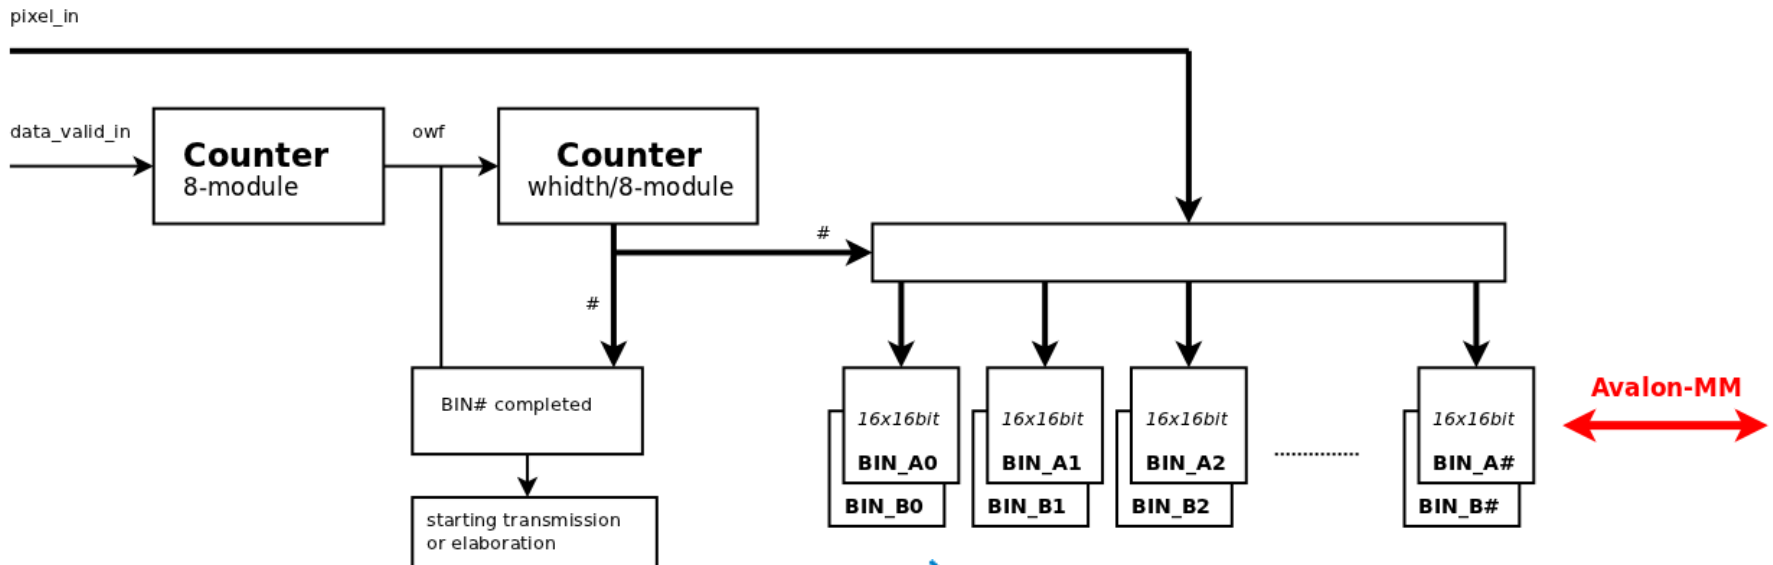
Performs the histogram extraction over a **8x8 pixel** cell from the previously extracted Gradient-frame



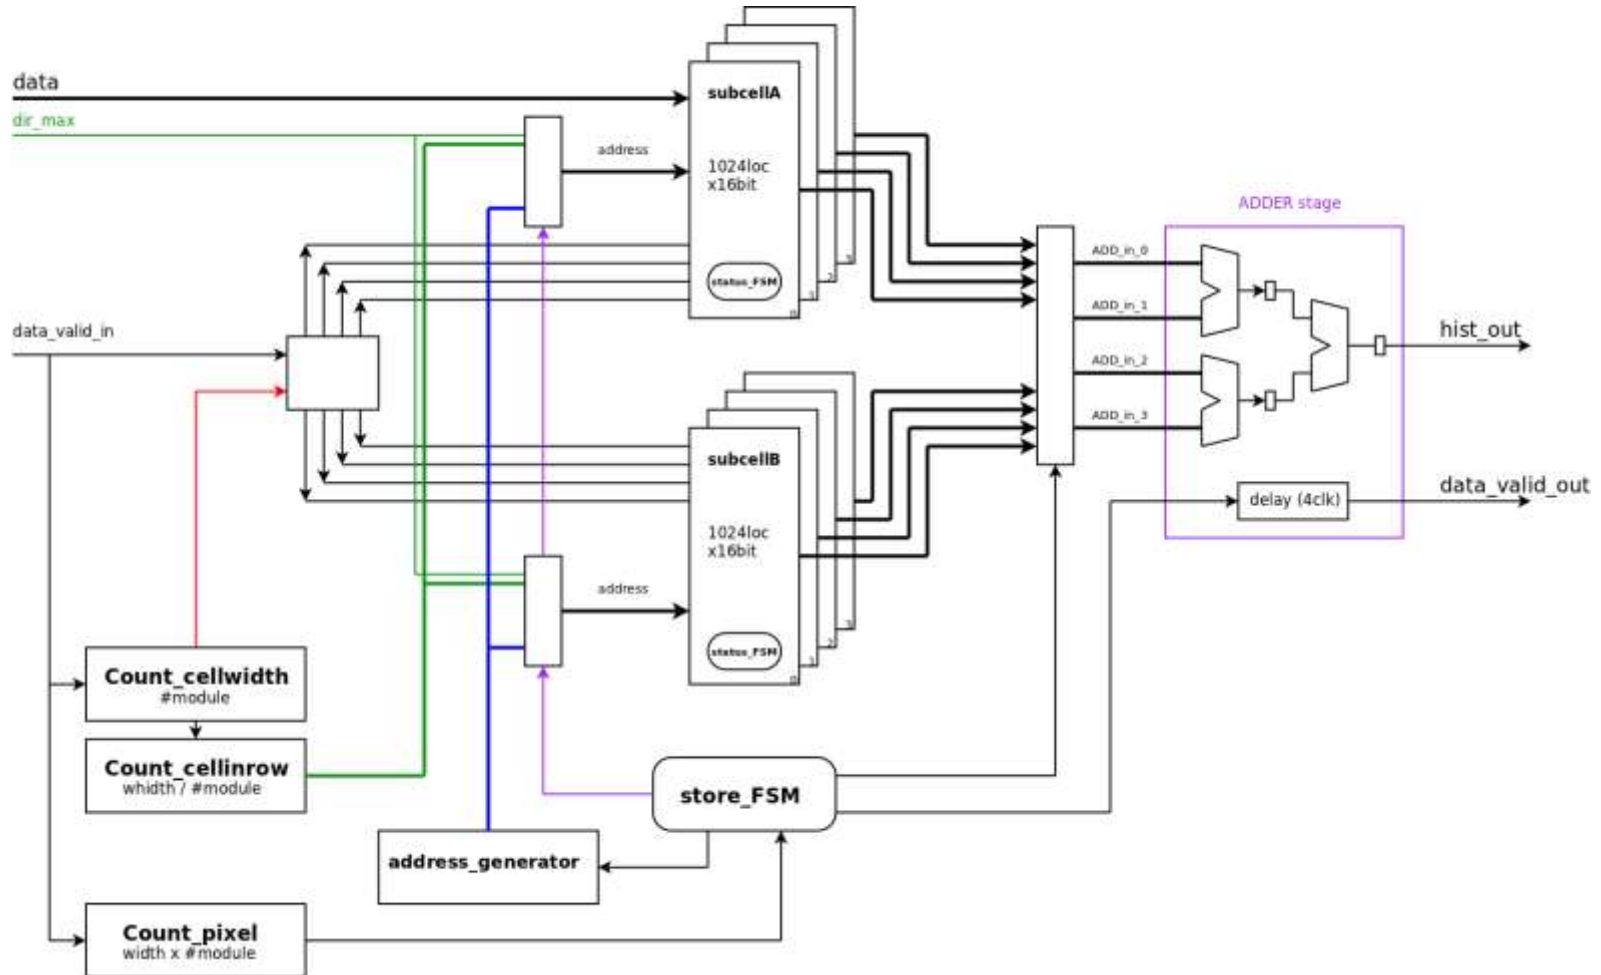
HistogramHW (2/3)

The HistogramHW is developed as a FSM able to manage a continuous stream of pixels

- performs a **single clock** read-modify-write operation (using parallel module)
- **constant data output latency** (depends on both the cell size and the image size)
 - about 2560 clock cycles @ Q-VGA and 8x8 cells

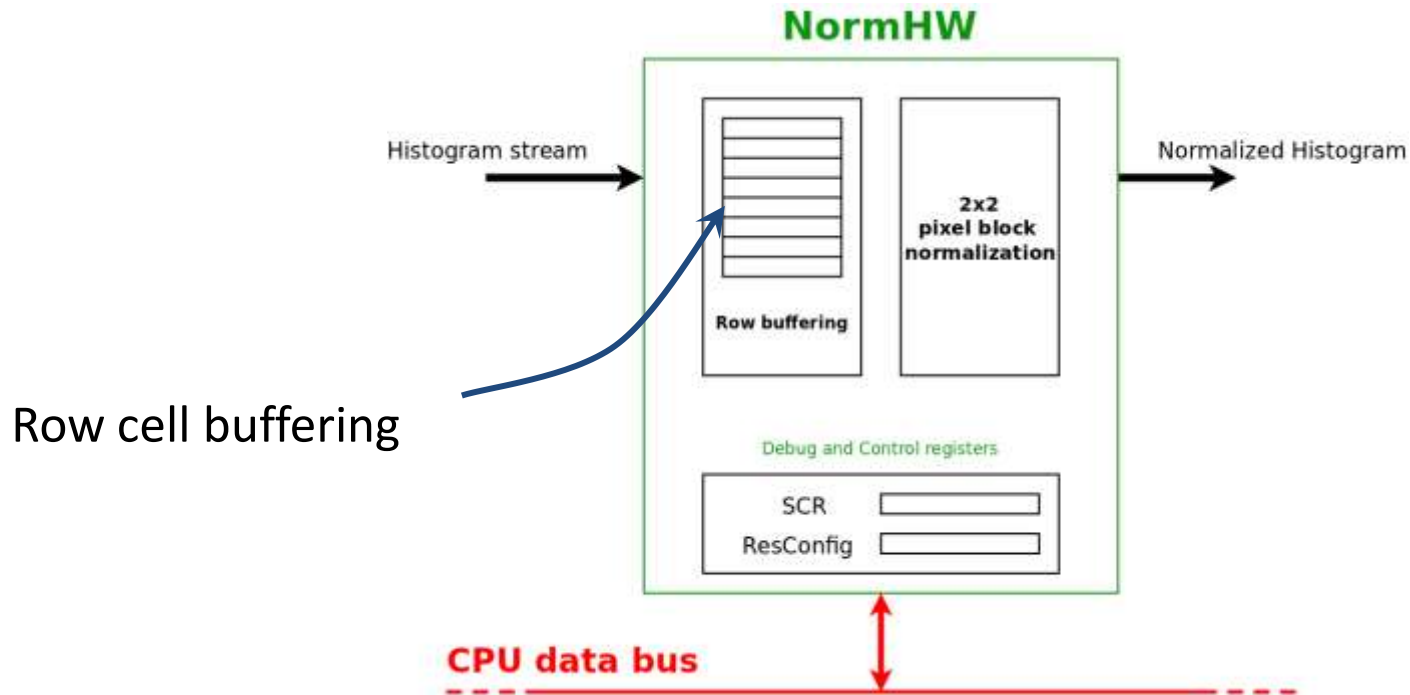


HistogramHW (3/3)



NormHW

Equalisation the luminance among close cells
(composed by a 2x2 cell)



Module still in development

Implementation results

FPGA resource use

VideoSampler

200 LUT

512 Byte (FIFO buffering)

GradientHW

1200 LUT

960 Byte (Row buffering)

32 DSP module 9x9

HistogramHW

850 LUT

16 kByte (Cell buffering)

NormHW

400 LUT

2 kByte (Block buffering)

module still in development

HOG implementation test

- Image captured from an OmniVision CMOS Camera OV9650
- Elaboration performed through the configurable HOG pipeline
- PC-interface use Altera *JtagAtlantic* interface to receive the datastream (USB link)



Processing latency

All the blocks inside the pipeline are implemented using the **streaming paradigm**:

- constant latency: the maximum value between all the block latencies
 - Maximum latency = Histogram latency (2560 clock cycles)
- works at the same fps as input
 - the maximum manageable frame rate depends on technological constraints
 - contingent case: 12 fps (~83ms)

Conclusions

- The HOG pipeline is implemented on FPGA using Hardware-Software codesign approach
- A new edge detector kernel is proposed in order to compute directly both the magnitude and the orientation of a vector, exploiting the hardware parallelism capability
- Hardware blocks implemented using the streaming paradigm:
 - edge detector: 2 clock cycles of latency
 - histogram: the latency depends on image size and cell size
 - Contingent case: Q-VGA and 8x8 cells, about 2560 clock cycles
 - total latency: the maximum value between the blocks latencies
 - Contingent case: histogram latency, about 2560 clock cycles

FPGA implementations of Histograms of Oriented Gradients for vehicle detection

Domain Specific Language Implementation

C. Bourrasset, J. Sérot and F. Berry

Institut Pascal- D.R.E.A.M - Aubière, France



Domain Specific Language approach


- High level abstraction programming
- Rapid prototyping
- FPGA programming becomes easy !
- HDL generated code is less efficiently than handwritten code

The CAPH programming language

- Based upon the **dataflow** model of computation
 - offers an **bridging** MoC between high-level, programmers oriented specification and low-level, HDL-based descriptions
- Well suited to **stream-processing** applications (operating *on-the-fly* on continuous streams of data coming directly from sensors)

Dataflow model

- An application = a collection of computing units (*actors*) exchanging *tokens* through unidirectional, *buffered* links (*FIFOs*)
- For each actor, a set of *firing rules* specifies when it consumes input tokens and produces output tokens



CAPH language

- Dataflow model need **actors** and **connection** between actors :
 - One sub-language for actor behavior description
 - One sub-language for network description

Actor description

```
type pixel= unsigned<8> ;
```

```
actor inv ()  
  in (a: pixel dc)  
  out (c: pixel dc )
```

```
rules
```

```
| a: '<' -> c: '<'  
| a: 'p' -> c: '255-p'  
| a: '>' -> c: '>'  
;
```

Network description

```
stream i: pixel dc from "dev:cam0";
```

```
stream o: pixel dc to "dev:mon1";
```

```
net o = inv i;
```



CAPH tokens

- A token in CAPH may represent
 - Datas (integers, boolean, floats)
 - Control signals ‘<’, ‘>’
- Example image



Gray level representation

103	45	50	...
60	34	24	...
150	42	210	...
...

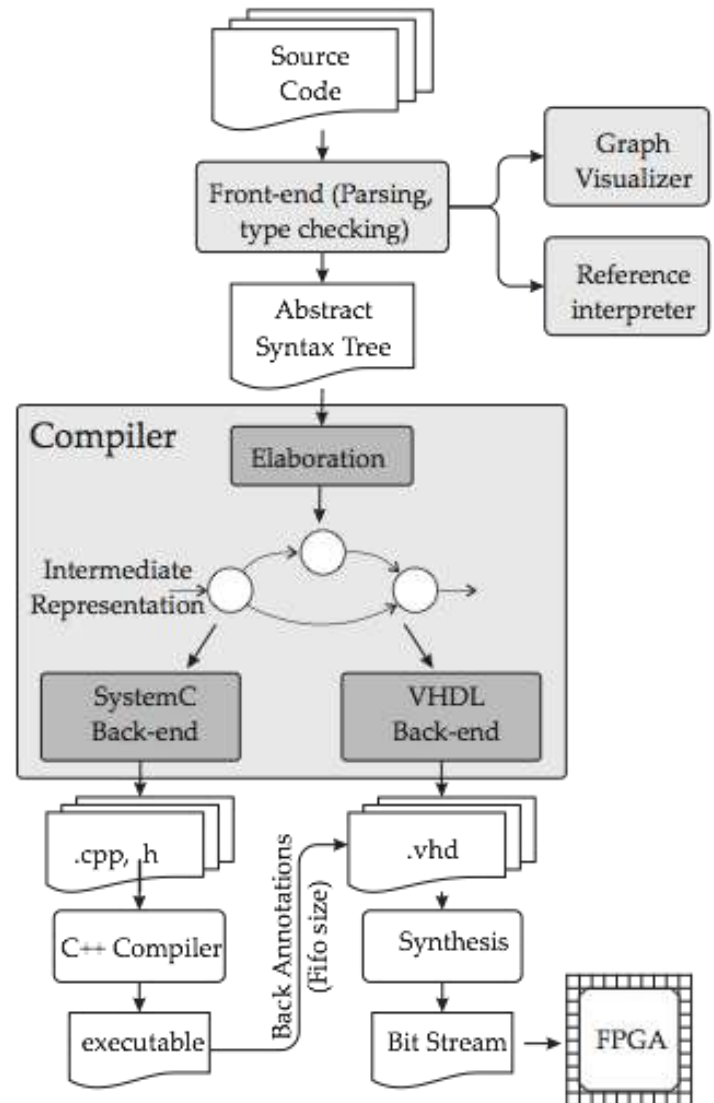
- CAPH tokens transcription

< < 103 45 50 ... > < 60 34 32 ... > < 150 40 210 .. > < ... > >

Start of Frame Start of Line End of Line End of Frame


Caph toolset

- **Graph visualizer** : .dot format
- **Reference interpreter** :
 - based on the fully formalized semantics
 - tracing, profiling and debugging
- **Compiler** :
 - elaboration of a target-independent IR
 - specialized backends ([SystemC](#), [VHDL](#))



HOG

Window size 320 x 240



- Gradient computation

$$|G| = |G_x| + |G_y|$$


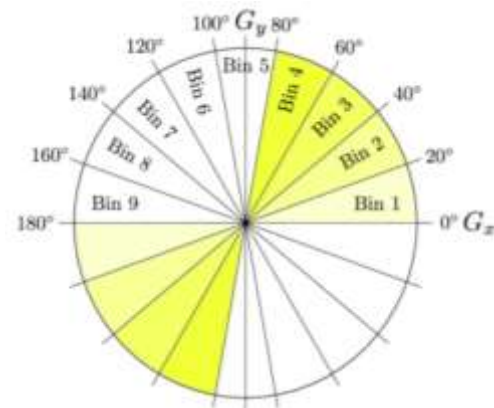
$$G_x = M_x * I \quad M_x = \begin{bmatrix} 1 & 0 & 1 \end{bmatrix}$$

$$G_y = M_y * I \quad M_y = \begin{bmatrix} 1 & 0 & 1 \end{bmatrix}^T$$

- Gradient orientation


$$G = \tan^{-1} \left(\frac{G_y}{G_x} \right)$$

- Discretization into 9 spaced angular bin over 0°-180°
- Generate 9 magnitude-weighted bin images



HOG

Window size 320 x 240



Source Code

```

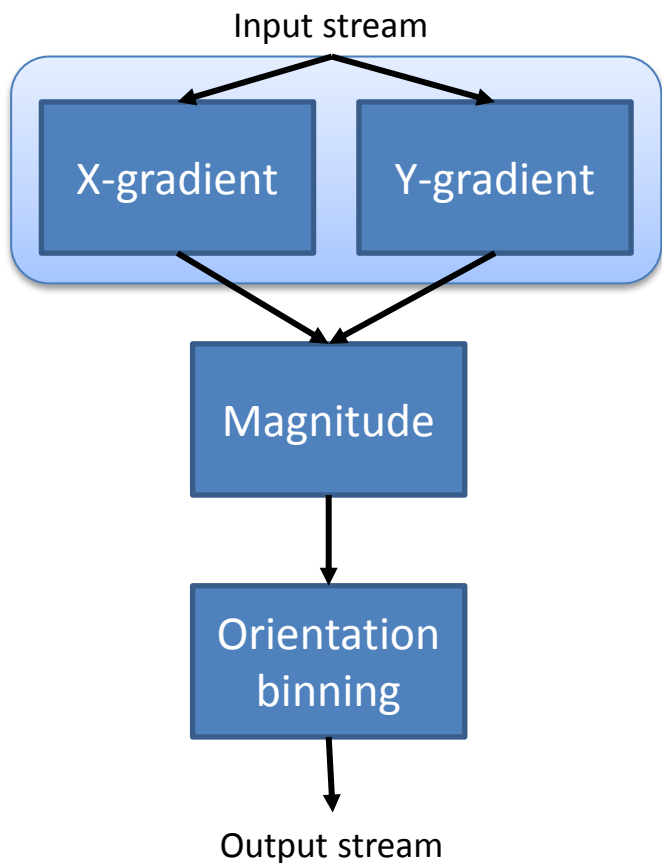
actor maddn(k:signed<9> array[3])
in (a:unsigned<8> dc,b:unsigned<8> dc, c:unsigned<8> dc)
out (o:signed<9>dc)
rules
  ( a,b,c) -> o
  |('<','<','<') -> '<'
  |('zpz','zp','p') -> '( zzp*k[0] + zp*k[1] + p*k[2])'
  |('>','>','>') -> '>';
  net dx = maddn [[1,0,(-1)]] (neigh2u13 i);
  net dy = maddn [[1,0,(-1)]] (neigh2u31r i);

  net neigh2u13 x =
    let y2 = d1p x in
    let y3 = d1p y2 in (x, y2, y3);
    
```

- d1p actor delays one pixel to the right


Example:

$\langle 1 \ 2 \ 3 \ 4 \rangle \xrightarrow{d1p} \langle 0 \ 1 \ 2 \ 3 \rangle$

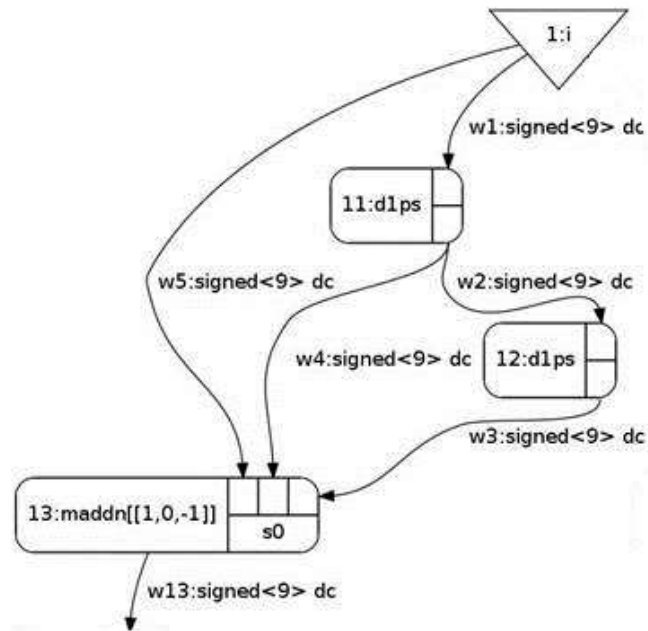
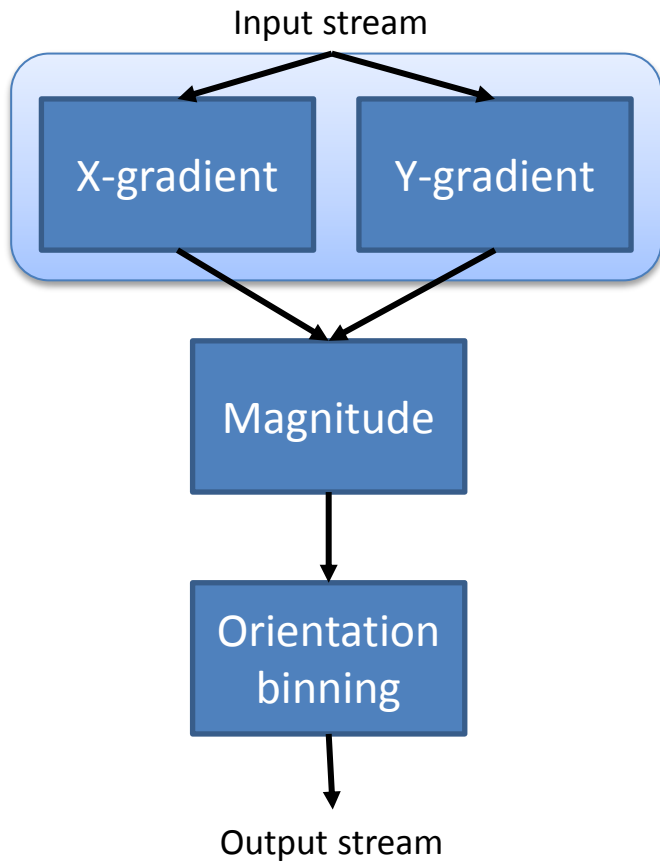


HOG

Window size 320 x 240




- Network of actors for x gradient

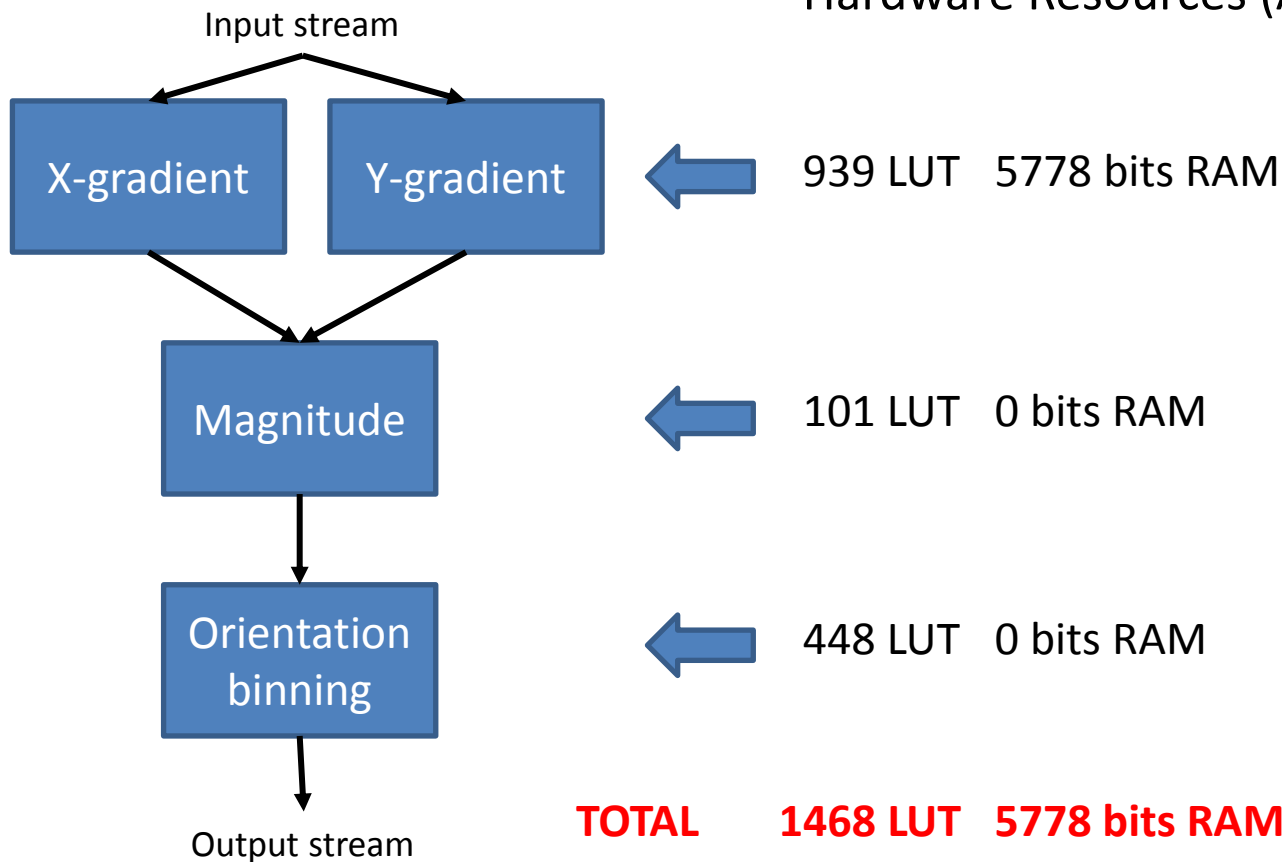


HOG


Window size 320 x 240




- Hardware Resources (Altera Cyclone III FPGA)



HOG



- Histogram extraction over a 8x8 pixel block (cell) from gradient frame




- 9 parallel streams computed by 9 actors in parallel
- Need to store 8 rows to compute histogram extraction (using FIFOs)


Hardware resources for histogram extraction

9500 LUT 7,8 kbits RAM

HOG



- Local contrast normalization
- Reduce the luminance variation
- L1-norm, $\mathbf{v} \rightarrow \mathbf{v} / (\|\mathbf{v}\|_1 + \epsilon)$



Hardware resources for normalization

900 LUT 2 kbits RAM

FPGA resource summary

Gradient	1468 LUT	5,7 Kbits
Histogram	9500 LUT	7,8 Kbits
Normalization	900 LUT	2 Kbits

- Histogram step can be optimized
- Generated HDL code synthetize LUT locks to memory job



CAPH Project

- Continuous development since 2009 (Major release soon)
- Substantial increase in abstraction level compared to classical HDLs
 - large gain in development times (x5-x10)
 - without significant performance penalty (< 30%)
- Several realistic applications implemented
 - motion detection, MPEG encoding, connected component labeling, Harris-Stephen POI detection,
- Toolset and manual freely available at

<http://dream.univ-bpclermont.fr/index.php/softmenu/caph>

Conclusion

VideoSampler

Co-design approach

200 LUT

512 Byte (FIFO buffering)

DSL approach

GradientHW

1200 LUT

960 Byte (Row buffering)

32 DSP module 9x9

1468 LUT

5.7 kbit

HistogramHW

850 LUT

16 kByte (Cell buffering)

9500 LUT

7,8 kByte

NormHW

400 LUT

2 kByte (Block buffering)

module still in development

600 LUT

2 kByte

210 Lines of Code



Thank you

